
EML Trees Are Universal Approximators

Anonymous Author(s)

Affiliation

Address

email

Abstract

1 The recently introduced EML (Exp-Minus-Log) function acts as continuous ana-
2 logue of NAND gates, providing a compositional building block capable of rep-
3 resenting elementary functions. In this work, we study the expressive power of
4 tree-structured compositions of EML functions. We show that such trees enjoy
5 a universal approximation property for functions in $W^{k,\infty}$ for $k \in \mathbb{N}$, drawing
6 on classical neural network approximation arguments while exploiting the abil-
7 ity to explicitly construct EML trees that mimic polynomial representations. We
8 further propose a learning algorithm for EML-type trees equipped with fitting
9 parameters, and demonstrate its feasibility in practical optimization problems. Our
10 results establish EML trees as a theoretically grounded framework for function
11 approximation.

12 1 Introduction

$$\text{EML}(x, y) = e^x - \text{Log } y \quad (1)$$

13 Analogy to NAND gates

14 Mention that EML is a mathematical curiosity, more than a new paradigm for function approximation

15 State that since we have all elementary functions, that we can recognize that we should expect an
16 approximation theorem, since we can build polynomial

17 2 Background

18 EQL, KAN, AI-Feynman, SINDy (/ PDEFind), Symbolic regression, SymTorch, Bacon (Langley)

19 Symbolic metamodel:

- 20 • Demystifying Black-box Models with Symbolic Metamodels (NeurIPS)

21 3 Universal Approximation Theorem

22 For our work, we consider a generalization of the EML function (1), denoted EML_{θ_i} , to include 6
23 learnable parameters each, and represented as

$$\text{EML}_{\theta_i}(x, y) = a_i e^{b_i x + c_i} + d_i \ln(e_i y + f_i), \quad (2)$$

24 where i enumerates the instance of the EML function in the tree and

$$\theta_i = \{a_i, b_i, c_i, d_i, e_i, f_i\}. \quad (3)$$




Figure 1: Sketch of the proof of the universal approximation theorem. **THIS IS AN OLDER SCHEMATIC; NEEDS UPDATE.**

For ..., we will often refer to individual EML_{θ_i} as “atoms” and the reusable functions or operations built using a collection of EML’s with fixed parameters as “blocks”.

Of course, one observes that setting $a_i = b_i = e_i = 1$, $d_i = -1$, and $c_i = f_i = 0$ recovers the original EML function (1).

The purposes behind this generalization are both technical and practical.

As will be detailed in the upcoming subsections, proving the approximation theorem relies on exactly representing a polynomial as an EML tree, but the technical lemma 2, which allows us to approximate the original function by a polynomial does not give us an estimate or an upper bound on the coefficients of these polynomials; thus, with the current framework, it is impossible to know the size of the EML tree that represents such polynomials, because in the original function (1), constants need to build recursively from smaller ones (starting from 1), so reaching coefficients badly depends on their magnitude.

Additionally, practically, as we will witness when we construct operations using these generalized EML functions, they are able to represent operations and functions more compactly, as will be demonstrated for the basic binary operations $(+, -, \times, \div)$ and others used in building polynomials.

It also performs better.

Another advantage is that we now no longer need to consider the number 1 as a possible input to the EML atoms, thus reducing the inputs to the EML atoms as either the variables or other EML atoms.

3.1 Setting

3.2 Statement of the Theorem

Theorem 1 (Universal Approximation Theorem in $W^{k,\infty}$).

3.3 Outline of the Proof

3.4 Remarks

Remark 1. For non positive functions.

Remark 2. For recovering trigonometric functions.

50 **Remark 3.** *There is an inherent nonuniqueness associated with these generalized EMLs. (comment*
 51 *on addition block for example).*

52 **Remark 4.** *Corollary to any domain via affine transformation.*

53 **Remark 5.** *Removing the need for approximating indicator functions by adding ReLU function to*
 54 *definition of EML (?).*

55 4 Implementation

56 The constructive proof of Theorem 1 establishes that the parametric EML grammar can express,
 57 exactly, any element of an arbitrarily fine polynomial approximant to a $W^{k,\infty}$ function. It does not
 58 by itself address whether such a tree can be *recovered* from data by gradient-based optimisation.
 59 In contrast to Odrzywolek [2026], which targets *symbolic* recovery of vanilla-EML trees over the
 60 discrete grammar $\{1, x, \text{EML}(\cdot, \cdot)\}$ and reports rapid collapse of recovery rate at depth ≥ 5 , we
 61 exploit the continuous parameters of the generalised atom $\text{EML}_{\theta_i}(x, y)$ to make the optimisation
 62 landscape amenable to first-order methods. This section reports a small, focused empirical study in
 63 support of the theorem’s practical content.

64 4.1 Setup

65 **Architecture.** We use the PEM tree of Section 1 directly: a complete binary tree T_θ of $2^d - 1$
 66 atoms, each carrying six free parameters $\theta_i = \{a_i, b_i, c_i, d_i, e_i, f_i\}$ and computing $\text{EML}_{\theta_i}(x, y)$.
 67 Children are fed *raw* between atoms, with the input variable x entering at the leaf-parents directly.
 68 For univariate x the parameter count is $\#\theta(d) = 6(2^d - 1) \in \{42, 90, 186\}$ at depths $d \in \{3, 4, 5\}$.

69 **Two parameterisations.** We consider two regimes:

70 **Real (R).** All $\theta_i \in \mathbb{R}^6$. To keep the operator real on $\mathbb{R}$, the logarithm’s argument is routed through
 71 $\ln(\text{softplus}(\cdot) + \epsilon)$ with $\epsilon = 10^{-6}$, a smoothed-positive surrogate for the principal-branch
 72 logarithm. This blocks the trigonometric path of the corresponding remark in Section 1 but
 73 eliminates branch-cut pathologies.

74 **Complex (C).** All $\theta_i \in \mathbb{C}^6$, stored as twin real Parameters (θ_i^R, θ_i^I) . Forward in $\mathbb{C}^{128}$ with the
 75 principal branch of $\ln$; the loss is computed on $\text{Re}(T_\theta(x))$. We add a small auxiliary penalty
 76 $\lambda \cdot |\text{Im } T_\theta|^2$ with $\lambda = 10^{-3}$ so that the trained predictor is honestly real-valued. Doubling the
 77 parameter count is the cost; gaining the trigonometric path is the intended benefit.

78 **Optimisation procedure.** We minimise mean-squared error on a uniform grid of $N = 200$ samples
 79 per target with the following two-stage strategy, applied per (target, depth, parameterisation) cell at a
 80 fixed wall-clock budget of 30 seconds:

- 81 1. **Adam multi-restart.** From independent Gaussian-initialised seeds, run 1500 steps of Adam with
 82 learning rate 5×10^{-3} , cosine-annealed to zero, gradient ℓ_2 -clip at 5.0. Track the best mean-squared
 83 error across restarts.
- 84 2. **L-BFGS refinement.** Apply L-BFGS with strong-Wolfe line search (≤ 200 iterations, history 20)
 85 to the best-of-restarts model and take the smaller of the pre- and post-refinement losses.

86 The number of Adam restarts per cell is determined by the budget rather than fixed and is reported in
 87 Table 1.

88 **Assumptions.** The reported numbers depend on three explicit choices. (i) Under R the operator is
 89 the softplus-stabilised log surrogate, not vanilla EML; the two coincide on positive arguments and
 90 diverge smoothly near zero. (ii) All targets are noise-free, dense samples on the stated domain; no
 91 claim is made about generalisation to sparse or noisy data. (iii) The wall-clock budget is small (30 s on
 92 a single CPU) and held fixed across cells; reported errors at the larger depths are optimisation-limited,
 93 not capacity-limited.

94 4.2 Targets and metric

95 We report results on five 1D targets, each chosen to align with a specific construction in the proof:

Table 1: Best relative RMSE per cell, PEM trees with real (R) and complex (C) parameters at depths 3, 4, 5. Strategy: Adam multi-restart + L-BFGS, 30 s wall-clock per cell. Parenthesised integer is the number of completed Adam restarts within the budget; parameter count is $\#\theta = 6(2^d - 1)$ for R and twice that for C. Bolded entry per row is the best of the six.

Target	Real (R), $\#\theta = 42, 90, 186$			Complex (C), $\#\theta = 84, 180, 372$		
	$d=3$	$d=4$	$d=5$	$d=3$	$d=4$	$d=5$
$x^3 - x$	$7.58 \cdot 10^{-3}$ (15)	$1.47 \cdot 10^{-3}$ (8)	$4.32 \cdot 10^{-3}$ (4)	$2.17 \cdot 10^{-2}$ (7)	$1.37 \cdot 10^{-2}$ (4)	$5.20 \cdot 10^{-2}$ (2)
$\tanh(2x)$	$4.20 \cdot 10^{-3}$ (15)	$1.22 \cdot 10^{-3}$ (8)	$2.10 \cdot 10^{-3}$ (4)	$1.77 \cdot 10^{-3}$ (7)	$5.16 \cdot 10^{-3}$ (4)	$1.80 \cdot 10^{-2}$ (2)
$\sin(x)$	$3.02 \cdot 10^{-3}$ (15)	$7.67 \cdot 10^{-3}$ (8)	$1.01 \cdot 10^{-2}$ (4)	$1.74 \cdot 10^{-2}$ (7)	$3.49 \cdot 10^{-3}$ (4)	$2.07 \cdot 10^{-2}$ (2)
e^{-x^2}	$4.13 \cdot 10^{-3}$ (15)	$1.12 \cdot 10^{-3}$ (7)	$4.66 \cdot 10^{-3}$ (4)	$2.20 \cdot 10^{-2}$ (7)	$5.27 \cdot 10^{-3}$ (3)	$1.27 \cdot 10^{-2}$ (2)
$\sin(3x) e^{-x^2/2}$	$3.23 \cdot 10^{-2}$ (15)	$9.80 \cdot 10^{-2}$ (7)	$9.09 \cdot 10^{-2}$ (4)	$6.45 \cdot 10^{-2}$ (6)	$3.91 \cdot 10^{-2}$ (3)	$3.57 \cdot 10^{-2}$ (2)




Figure 2: Best relative RMSE vs PEM tree depth, five targets, real (solid) and complex (dashed) parameterisations, 30 s wall-clock per cell. The dotted line marks the 1% relRMSE threshold. Real PEM reaches sub-percent error on 4/5 targets by depth 4; complex PEM is competitive only on $\sin(3x) e^{-x^2/2}$, where its larger capacity helps despite the halved restart count.

- $x^3 - x$ on $[-1.5, 1.5]$: a degree-3 polynomial; the local approximant of Lemma 2 in its simplest non-trivial form. The constructive bound is $5k - 3 = 12$ atoms for degree $k = 3$.
- $\tanh(2x)$ on $[-2, 2]$: hyperbolic tangent, used in the partition-of-unity construction (Section 6); the explicit EML realisation is in Figure 5.
- $\sin(x)$ on $[-\pi, \pi]$: the trigonometric case highlighted in the corresponding remark; the vanilla-EML construction routes through $\ln(-1) = i\pi$, unreachable from R.
- e^{-x^2} on $[-2.5, 2.5]$: smooth analytic, no oscillation — baseline for “easy” targets.
- $\sin(3x) e^{-x^2/2}$ on $[-3, 3]$: oscillation $\times$ envelope; the hardest target in this set.

We report *relative RMSE*, $\text{relRMSE} = \|T_\theta(x) - f(x)\|_2 / \|f(x)\|_2$, on the training grid, best-of-restarts.

4.3 Results

Table 1 reports the best relRMSE per cell. Figure 2 plots relRMSE against depth on a log scale.

What the numbers say. Three observations support the theorem’s practical content. (i) Real-parameterised PEM trees achieve sub-percent relRMSE on 4 of 5 targets at depth 4 within 30 s of

single-CPU compute, including the polynomial target $x^3 - x$ which falls to 1.47×10^{-3} . This is a direct empirical reflection of the constructive route of the theorem (polynomial approximation, then exact EML representation): gradient methods reach close to the constructed optimum, even though the construction itself is not used as an initialisation. (ii) The trigonometric target $\sin(x)$ does *not* require the complex parameterisation in our setting: real PEM at depth 3 already reaches 3.0×10^{-3} . The softplus-stabilised log is expressive enough on this domain that the complex path is not the binding constraint at the budget shown. (iii) The hardest target, $\sin(3x)e^{-x^2/2}$, is the one clear case where the complex parameterisation pays off: 3.6×10^{-2} at depth 5 (C) versus 9.1×10^{-2} at the same depth (R). The composite oscillation-envelope target rewards the extra capacity even though C completes only half as many restarts in the same wall-clock budget.

The general pattern visible in Figure 2 — error falling sharply from depth 3 to depth 4, then rising at depth 5 for several real targets — is a wall-clock-budget effect: deeper trees have more parameters but, at fixed budget, fewer restarts to find good basins. With unlimited compute this artefact would disappear; under fixed compute it suggests an “elbow” depth around $d = 4$ for these 1D targets.

Computational cost. The full sweep (5 targets $\times$ 3 depths $\times$ 2 parameterisations = 30 cells) completes in approximately 12 minutes on a single laptop CPU (Apple M-series, no GPU). Per-cell wall-clock is 30 s by construction; the number of completed Adam restarts ranges from ~ 15 at ($d=3, R$) down to 2 at ($d=5, C$). No GPU acceleration or parallelism was used.

4.4 Inspection of the fitted trees: are they symbolic?

The PEM atom carries the parameter-free EML as the corner case $\theta = (1, 1, 0, -1, 1, 0)$, and the constructive proof in Sections D–E realises every elementary function as a tree whose every atom has $\theta_i \in \{0, \pm 1\}^6$. A natural question is whether the fitted trees recover this structure, or whether they exploit the continuous $\mathbb{R}^6$ freedom of each atom to find function-level fits that bear no relation to the symbolic blocks.

Procedure. For each of the five targets, we refit a depth-3 PEM (7 atoms, 42 parameters) with 24 Adam restarts plus an L-BFGS refinement. We then greedily snap each parameter to the nearest member of $\{-2, -1, 0, 1, 2\}$, accepting a snap if the loss does not grow by more than 30%, and refit the un-snapped parameters with the snapped ones held fixed. Atoms whose six parameters all snap are reported as *fully discrete*.

Result. Table 2 summarises. Across all five targets, *at most 7/42 (17%) parameters are snappable, and zero atoms become fully discrete*. On $\tanh(2x)$ and e^{-x^2} exactly 0 parameters snap.

Table 2: Greedy snap of fitted depth-3 PEM parameters to $\{-2, -1, 0, 1, 2\}$ at 30% relative-loss tolerance.

Target	relRMSE pre	relRMSE post	snapped	fully-discrete atoms
$x^3 - x$	$7.3 \cdot 10^{-3}$	$6.9 \cdot 10^{-3}$	7/42 (17%)	0
$\tanh(2x)$	$3.8 \cdot 10^{-3}$	$3.5 \cdot 10^{-3}$	0/42 (0%)	0
$\sin(x)$	$1.9 \cdot 10^{-3}$	$1.7 \cdot 10^{-3}$	3/42 (7%)	0
e^{-x^2}	$3.9 \cdot 10^{-3}$	$3.2 \cdot 10^{-3}$	0/42 (0%)	0
$\sin(3x)e^{-x^2/2}$	$1.18 \cdot 10^{-1}$	$7.7 \cdot 10^{-2}$	3/42 (7%)	0

Example. For $x^3 - x$ (the most snappable target), the fitted parameter table is

atom i	a_i	b_i	c_i	d_i	e_i	f_i
root	1.516	-0.654	0.796	-2.080	0.652	0.029
left child	1.218	-0.436	0.147	-0.755	-0.034	1.501
right child	1.803	0.473	0.138	-2.348	0.115	-0.785
leaf-parent 1	1.299	-1.359	0.847	-0.143	0.127	0.449
leaf-parent 2	1.272	-2.646	0.116	-0.065	0.357	0.841
leaf-parent 3	1.267	-0.072	0.504	-1.190	0.114	0.471
leaf-parent 4	0.913	2.388	0.475	-0.672	-0.219	1.080

143 The a column clusters around 1–2 but contains no exact integers; the d column ranges over -2.35
 144 to -0.07 with no values near the original-EML default $d = -1$; the b column spans -2.65 to 2.39
 145 with no recognisable structure. The constructive proof builds $x^3 - x$ from $5k - 3 = 12$ atoms with
 146 $a, d \in \{0, \pm 1\}$ exactly and the coefficients absorbed into c via $c = \log m$. The fitted 7-atom tree
 147 matches the function value to 0.7% but does so by additive interference of stacked exp/log curves
 148 with arbitrary continuous slopes, not by anything resembling the proof’s polynomial expansion.

149 **Why.** A PEM atom has $\mathbb{R}^6$ of freedom; the symbolic constructions occupy a measure-zero subset.
 150 From Gaussian initialisation, gradient descent almost surely lands somewhere else. The loss landscape
 151 is shape-determined: many parameter settings produce indistinguishable L^2 values on the training
 152 grid, and the optimiser settles on whichever continuous-coefficient minimum its trajectory reaches
 153 first.

154 **Implication.** The §4.1–4.5 results support the *approximation* content of the theorem: trees reach
 155 small error in practice. They do *not* support an *interpretability* claim about the recovered trees.
 156 Recovering discrete-parameter trees with the symbolic structure of Sections D–E would require
 157 additional structural regularisation (sparsity penalty toward $\{0, \pm 1\}$, warm-start from the construction,
 158 or a two-stage train-then-snap-then-refit at tighter tolerance). We leave these for future work.

159 4.5 Limitations

160 We list the limitations of this empirical study honestly, in priority order.

- 161 • **Single-run-of-experiment.** Tabled values are the best loss from one run of the multi-restart
 162 procedure. Variance across full reruns is not reported. Order-of-magnitude conclusions are robust;
 163 comparisons within a factor of ~ 2 should not be over-interpreted.
- 164 • **Real-softplus surrogate.** The R parameterisation is not vanilla EML. Trees fit under R coin-
 165 cide with EML on the positive-argument slice but smoothly diverge near zero. The complex
 166 parameterisation uses the operator literally.
- 167 • **Univariate, smooth, dense, noise-free.** Multivariate fitting, sparse sampling, and noise robustness
 168 are all open and likely harder; the theorem is dimension-agnostic but our experiments are not.
- 169 • **Fixed-budget effects on depth comparison.** Deeper trees receive fewer restarts at fixed wall-clock,
 170 so the depth-vs-error curve confounds capacity gain with restart-count loss. The dip at depth 4 for
 171 several targets is the visible signature.
- 172 • **Optimisation, not approximation, theory.** The reported numbers are achieved errors, not bounds.
 173 They are upper-bounded by the theorem; whether they meet the rate the theorem predicts (and at
 174 what budget the gap closes) is a question we do not address here.
- 175 • **Function-level fits, not symbolic recoveries.** Section 4.4 shows that the fitted trees are continuous-
 176 coefficient approximations, not symbolic constructions matching the proof’s blocks. Recovering
 177 interpretable trees is open (see that section).

178 5 Discussion

- 179 • Generalization to other spaces (Besov)
- 180 • Remark on analytic functions
- 181 • Finding a way to use the vanilla EMLs by following another proof technique that involves
 182 finding the best polynomial approximant with bounded coefficients and finding upper bounds
 183 for constructing these constants (which is a challenging task to get the most optimal upper
 184 bounds for getting high constants).
- 185 • We could maybe make due with less than 6 parameters per EML atom.
- 186 • Trying to find better techniques to optimize these trees.
- 187 • Probably shorter trees exist that do the same thing; but, in this case, they were constructed
 188 like this to be human readable and easy to verify and interpretable.
- 189 • Can we find another function that trades depth for “width“, by resorting to making the atoms
 190 take in more inputs (similar to network constructions)

- 191 • Practically, the logarithm takes in only positive values, so we could put an absolute value
 - 192 inside the logarithm to circumvent any potential issue with domain
 - 193 • Discussion when talking about possible use for interpretable scientific discovery (Wyder
 - 194 et al. [2025], Yermakov et al. [2025])
- 195 **Open Question 1.** *Can we*

A Notation

$\mathbb{N}^d$

$\times$ denotes cartesian product

B Sobolev Spaces

We recall the definitions of some basic functional spaces.

The Lebesgue spaces, denoted $L^p(\Omega)$ are defined as

$$L^p(\Omega) = \left\{ f : \Omega \rightarrow \mathbb{R} \mid \int_{\Omega} |f(x)|^p dx < \infty \right\}, \quad (4)$$

for $1 \leq p < \infty$, equipped with the norm

$$\|f\|_{L^p(\Omega)} = \left(\int_{\Omega} |f(x)|^p dx \right)^{\frac{1}{p}}. \quad (5)$$

When $p = \infty$, where $L^\infty(\Omega)$ consists of functions that are essentially bounded and whose norm is

$$\|f\|_{L^\infty(\Omega)} = \operatorname{ess\,sup}_{x \in \Omega} |f(x)| = \inf \{ C \geq 0 \mid |f(x)| \leq C \text{ for almost every } x \in \Omega \}. \quad (6)$$

Moreover, the Sobolev space $W^{k,p}$ is defined as

$$W^{k,p}(\Omega) = \{ f \in L^p(\Omega) \mid D^\alpha f \in L^p(\Omega) \text{ for all } \alpha \in \mathbb{N}_0^d \text{ with } |\alpha| \leq k \}. \quad (7)$$

We also define associated seminorms. For $p < \infty$, we have

$$|f|_{W^{m,p}(\Omega)} = \left(\sum_{|\alpha|=m} \|D^\alpha f\|_{L^p(\Omega)}^p \right)^{\frac{1}{p}}, \quad \text{for } m = 0, \dots, k, \quad (8)$$

and for $p = \infty$, we have

$$|f|_{W^{m,\infty}(\Omega)} = \max_{|\alpha|=m} \|D^\alpha f\|_{L^\infty(\Omega)}, \quad \text{for } m = 0, \dots, k. \quad (9)$$

Based on these seminorms, we can define the following norms for $p < \infty$,

$$\|f\|_{W^{k,p}(\Omega)} = \left(\sum_{m=0}^k |f|_{W^{m,p}(\Omega)}^p \right)^{\frac{1}{p}}, \quad (10)$$

and for $p = \infty$,

$$\|f\|_{W^{k,\infty}(\Omega)} = \max_{0 \leq m \leq k} |f|_{W^{m,\infty}(\Omega)}. \quad (11)$$

In particular, when $p = 2$, the Sobolev spaces $H^k(\Omega) = W^{k,2}(\Omega)$ (with $k \in \mathbb{N}_0$) are Hilbert spaces with the corresponding norms $\|\cdot\|_{H^k(\Omega)} = \|\cdot\|_{W^{k,2}(\Omega)}$ and seminorms $|\cdot|_{H^m} = |\cdot|_{W^{m,2}}$ for any integer m such that $0 \leq m \leq k$.

Additionally, the space of k -times continuously differentiable functions $C^k(\Omega)$, with the associated norm $\|f\|_{C^k(\Omega)} = \|f\|_{W^{k,\infty}(\Omega)}$.

We will make frequent use of the following Banach algebra property.

Lemma 1. *Let $d \in \mathbb{N}$, $k \in \mathbb{N}_0$, and $\Omega \subset \mathbb{R}^d$. If $f, g \in W^{k,\infty}(\Omega)$, then*

$$\|fg\|_{W^{k,\infty}(\Omega)} \leq 2^k \|f\|_{W^{k,\infty}(\Omega)} \|g\|_{W^{k,\infty}(\Omega)}. \quad (12)$$

C Useful Definitions and Lemmas

We recall the definition of a star-shaped set, which is a necessary geometric requirement for the approximation of functions by polynomials in Sobolev spaces (in Lemma 2).

219 **Definition 1.** A set $A \subset \mathbb{R}^d$ is star-shaped with respect to every point in a subset $B \subset A$ if

$$\forall y \in B, \forall x \in A, [y, x] \subset A, \quad (13)$$

220 where $[y, x] = \{(1-t)y + tx \mid t \in [0, 1]\}$ is the straight line in $\mathbb{R}^d$ connecting the point y to the point
221 x .

222 We recall the Bramble-Hilbert lemma for Sobolev spaces, as proved in De Ryck et al. [2021], with
223 similar statements appearing in Durán [1983] and Verfürth [1999].

224 **Lemma 2.** Let $\Omega \subset \mathbb{R}^d$ be an open and bounded set of diameter $0 < h < e^{-1/2}d^{-3/2}$ which is
225 star-shaped with respect to every point in an open ball $B \subset \Omega$ with diameter ph . Then, for every
226 $f \in W^{s,\infty}(\Omega)$, there exists a polynomial $\hat{f}$ of degree at most $s-1$ such that for any $k \in \mathbb{N}_0$ with $k < s$,
227 it holds that

$$\|f - \hat{f}\|_{W^{k,\infty}(\Omega)} \leq \frac{\pi^{1/4} \sqrt{s} (d\sqrt{de}h)^{s-k}}{(s-k-1)!} |f|_{W^{s,\infty}(\Omega)}. \quad (14)$$

228 D Required Building Blocks

229 In this section, we present the explicit constructions of the basic binary operations, exponentiation,
230 and the hyperbolic tangent function, all of which are used in the approximation proof (Theorem 1).

231 D.1 Basic Binary Operations

232 In Figure 3, we provide the EML trees (with their associated parameters) to produce the basic binary
233 operations of addition, subtraction, multiplication, and division. The addition and subtraction blocks
234 are used in the constructions of the multiplication and division blocks respectively.

235 As observed in Figure 3, we note that the size and depth requirements of these binary operations are
summarized in Table 3.

Operator	$N(\cdot)$	Depth($\cdot$)
$x +_\theta y$	3	2
$x -_\theta y$	3	2
$x \times_\theta y$	6	4
$x \div_\theta y$	6	4
x_θ^n, mx_θ^n	2	2
$\tanh_\theta(x), m \tanh_\theta(x)$	8	5

Table 3: Size and depth of the operations in Appendix D.

236

237 D.2 Exponentiation

238 In our proof, we will also make frequent use of exponentiation to a constant power. We provide
239 explicit constructions for EML trees that perform x^n and mx^n for constants $m, n \in \mathbb{R}$ in Figure 4.

240 Note that in Figure 4b, we construct mx^n directly, which we use to represent the monomial terms of
241 polynomials (in Appendix E). This is because it is more efficient to incorporate the coefficient within
242 the exponentiation block itself, rather than introducing a separate multiplication block to combine a
243 constant with the x^n block.

244 For both exponentiation blocks (4a and 4b), the size and depth of the associated EML trees are 2 (and
245 appended to Table 3).

246 D.3 Hyperbolic Tangent

247 The hyperbolic tangent is constructed, using the exponential definition, by first getting $e^{2x} - 1$ and
248 $e^{2x} + 1$ (via two separate EML blocks) and then dividing them, as in Figure 5. The size of the tanh
249 block is 8 and its depth is 5, which are appended to Table 3.




Figure 3: The Basic Binary Operations




Figure 4: Exponentiating

250 We can similarly obtain $m \tanh x$, which is useful for the construction of the approximate partitions
 251 of unity, by amending the division block to include the prefactor m . This is done by changing θ_3 in
 the division block (Figure 3d) to be $\theta_3 = \{m, 1, 0, 0, 1\}$.




Figure 5: Hyperbolic Tangent (8)

252

253 E Constructing Polynomials

254 The following two subsections detail how to construct univariate and multivariate polynomials using
 255 EML trees and provide upper bounds for the number of required EML blocks in each case.

256 E.1 Univariate Polynomials

257 A univariate polynomial of degree k , denoted by

$$p(x) = \sum_{i=0}^k a_i x^i = a_0 + a_1 x + a_2 x^2 + \cdots + a_k x^k, \quad (15)$$

258 can be (naively) constructed using EML blocks as in Figure 6, wherein we used two simple constituent
 259 blocks: addition (Figure 3a) and exponentiation to a constant with a prefactor (Figure 4b).




Figure 6: Univariate Polynomial (5s - 3)

260 The upper bound on the number of required EML blocks is obtained by noting that we used

- 261 • $k - 1$ addition blocks (each requiring 3 EML atoms), and
- 262 • k exponentiation with a prefactor blocks (each requiring 2 EML atoms).

263 So, constructing an EML tree p_θ such that $p_\theta = p$ (as per the above procedure) requires $N(T_\theta) =$
 264 $3(k - 1) + 2k = 5k - 3$ and $\text{Depth}(T_\theta) = 2 + 2(k - 1) = 2k$, since the longest path starts with the x
 265 feeding into the a_0 block and passes through all $k - 1$ addition blocks.

Hence, we phrase the above conclusion as the lemma:

Lemma 3 (Construction of Univariate Polynomials). *Given a univariate polynomial p of degree k , there exists an EML tree p_θ (following the explicit construction in Figure 6), such that $p_\theta = p$ with $N(p_\theta) \leq 5k - 3$ and $\text{Depth}(p_\theta) \leq 2k$.*

E.2 Multivariate Polynomials

We recall that, for $\mathbf{x} = (x_1, \dots, x_d) \in \mathbb{R}^d$, a multivariate polynomial p of degree k is of the form

$$p(\mathbf{x}) = \sum_{|\alpha| \leq k} c_\alpha x_1^{\alpha_1} x_2^{\alpha_2} \dots x_d^{\alpha_d}, \quad (16)$$

where $\alpha = (\alpha_1, \alpha_2, \dots, \alpha_d)$ is a multi-index with $|\alpha| = \alpha_1 + \alpha_2 + \dots + \alpha_d$, and c_α are the associated coefficients. Such a polynomial (of degree $\leq k$) has $\binom{d+k}{k}$ monomial terms.

To construct an EML tree for the polynomial (16), (1) we construct the monomials $c_\alpha x_1^{\alpha_1}, x_2^{\alpha_2}, \dots, x_d^{\alpha_d}$ for each α with $|\alpha| \leq k$, (2) multiply these monomials together (separately for each α), (3) do steps (1) and (2) for every monomial, and (4) add them together. A breakdown of the required EML tree sizes and depths for each step is as follows:

- (1) For a certain α , each monomial, of which there are d , requires an exponentiation block (and the first has a prefactor), totaling $2d$ EML atoms;
- (2) Again for each α , we need to multiply the d monomials, requiring $d - 1$ multiplication blocks and, thus, $6(d - 1)$ EML atoms; thus, each term, considering both getting monomials and multiplying them requires $2d + 6(d - 1) = 8d - 6$ EML atoms;
- (3) Since we have $\binom{d+k}{d}$ terms in the polynomial, we do steps (1) and (2) $\binom{d+k}{d}$ times, thus requiring $(8d - 6)\binom{d+k}{d}$ EML atoms;
- (4) We finally sum over all terms, requiring $\binom{d+k}{d} - 1$ addition blocks, each consisting of 3 EML atoms.

Thus, in total, constructing the polynomial p in (16) takes

$$(8d - 6)\binom{d+k}{d} + 3\left(\binom{d+k}{d} - 1\right) = (8d - 3)\binom{d+k}{d} - 3 \quad (17)$$

EML atoms, and whose depth can be calculated (by drawing a schematic like Figure 6) to be

$$2 + 4(d - 1) + 2\left(\binom{d+k}{d} - 1\right) = 4d + 2\binom{d+k}{d} - 4, \quad (18)$$

since the x 's feed into the first $c_\alpha x^{\alpha_1}$ block (of depth 2), pass through $d - 1$ multiplication (each of depth 4), and then through $\binom{d+k}{d} - 1$ addition blocks (each of depth 2).

The above justified the following lemma:

Lemma 4 (Construction of Multivariate Polynomials). *Given $p : \mathbb{R}^d \rightarrow \mathbb{R}$, a multivariate polynomial of degree k , there exists an EML tree p_θ , such that $p_\theta = p$ with*

$$N(p_\theta) \leq (8d - 3)\binom{d+k}{d} - 3, \quad \text{and} \quad \text{Depth}(p_\theta) \leq 4d + 2\binom{d+k}{d} - 4. \quad (19)$$

F Approximating the partition of unity

In this section, we aim to construct an approximate partition of unity, similar to Section 4 of De Ryck et al. [2021]. We do this because indicator functions are unsuitable to be represented by EML trees due to the fact that they contain a discontinuity and cannot be suitably represented by chaining elementary functions constructed from compositions of EML. For this reason, we construct smooth approximations of indicator functions on subsets of $[0, 1]^d$ via tanh functions. When suitably shifted and parametrized, these functions output values close to 1 within the subset and close to 0 outside of it.

302 F.1 Properties of the approximate partition of unity

303 We first recall some of the properties of the partition of unity of De Ryck et al. [2021].

304 For a point $y \in \mathbb{R}$ and a parameter $\alpha \in \mathbb{R}$ (to be chosen later), we define the one-dimensional
305 approximate indicator functions as

$$\begin{aligned}\rho_1^N(y) &= \frac{1}{2} - \frac{1}{2}\sigma\left(\alpha\left(y - \frac{1}{N}\right)\right), \\ \rho_m^N(y) &= \frac{1}{2}\sigma\left(\alpha\left(y - \frac{m-1}{N}\right)\right) - \frac{1}{2}\sigma\left(\alpha\left(y - \frac{m}{N}\right)\right) \quad \text{for } 2 \leq m \leq N-1, \\ \rho_N^N(y) &= \frac{1}{2}\sigma\left(\alpha\left(y - \frac{N-1}{N}\right)\right) + \frac{1}{2}.\end{aligned}\tag{20}$$

306 We can clearly see that if α is large enough, these act as soft indicator functions over the intervals
307 $\left(\frac{m-1}{N}, \frac{m}{N}\right)$ (depending on the subscript of ρ).

308 To choose α , we note that since $\sigma(x) \rightarrow 1$ as $x \rightarrow \infty$ and its derivatives also saturate ($\sigma^{(m)}(x) \rightarrow 0$
309 with $x \rightarrow \infty$), we can choose $R > 0$ large enough such that $|\sigma^{(m)}|$ is decreasing over $[R, \infty)$ for all
310 $1 \leq m \leq k$.

311 So, given $\varepsilon > 0$, we choose α , depending on N and ε , large enough such that

$$\frac{\alpha}{N} \geq R, \quad 1 - \sigma\left(\frac{\alpha}{N}\right) \leq \varepsilon, \quad \alpha^m \left| \sigma^{(m)}\left(\frac{\alpha}{N}\right) \right| \leq \varepsilon \text{ for all } 1 \leq m \leq k,\tag{21}$$

312 justified also by Lemma A.4 in De Ryck et al. [2021]. Consequently, α can be chosen to satisfy the
313 above conditions and has the explicit value

$$\alpha = N \max\left\{R, \ln\left(\frac{(2k)^{k+1}(Nk)^k}{e^k \varepsilon}\right)\right\},\tag{22}$$

314 (see Lemma A.5 in De Ryck et al. [2021] for justification).

315 Importantly, we define

$$\Phi_j^N(x) = \prod_{i=1}^d \rho_{j_i}^N(x_i),\tag{23}$$

316 which acts as our approximate partition of unity. We also define the set $\mathcal{V}_d$ as

$$\mathcal{V}_d = \left\{v \in \mathbb{Z}^d \mid \max_{1 \leq i \leq d} |v_i| \leq 1\right\},\tag{24}$$

317 which allow us choose only the cubes directly surrounding a cube of interest.

318 Lemmas 5 and 6 below prove that the set of functions $\{\Phi_j^N\}_{j \in \mathcal{A}^N}$ acts as an approximate partition of
319 unity in the sense that they satisfy the following two properties: for every $j \in \mathcal{A}^N$,

$$\sum_{v \in \mathcal{V}_d} \Phi_{j+v}^{N,d} \approx 1 \quad \text{and} \quad \sum_{\substack{v \notin \mathcal{V}_d, \\ j+v \in \mathcal{A}^N}} \Phi_{j+v}^{N,d} \approx 0.\tag{25}$$

320 We recall the below two lemmas from De Ryck et al. [2021]. The first provides a quantitative bound
321 for how close to 1 is the contribution of Φ_j^N and its direct neighbors on I_j^N .

322 **Lemma 5** (Lemma 4.1 of De Ryck et al. [2021]). *If $0 < \varepsilon < 1/4$, then*

$$\left\| \sum_{v \in \mathcal{V}_d} \Phi_{j+v}^N - 1 \right\|_{W^{k,\infty}(I_j^N)} \leq 2^{dk} d\varepsilon.\tag{26}$$

323 The second lemma guarantees that the contribution of the functions far away from the cube I_j^N is
324 small.

325 **Lemma 6** (Lemma 4.2 of De Ryck et al. [2021]). *Let $k \in \mathbb{N}_0$ and $v \in \mathbb{Z}^d$ with $v_i \geq 2$ for all $1 \leq i \leq N$.
326 Then, it holds that*

$$\|\Phi_{j+v}^N\|_{W^{k,\infty}(I_j^N)} \leq \max\{1, (2k)^{2k} \alpha^k\} \varepsilon.\tag{27}$$

327 F.2 Constructing the approximate partition of unity using EML trees

328 In Figure 7, we construct the ρ_m^N functions for $2 \leq m \leq N-1$, according to the definition in (20), and provide the associated parameters.




Figure 7: Construction of ρ_m^N for $2 \leq m \leq N-1$

329

330 This construction requires 4 EML blocks, 2 tanh blocks with $1/2$ constants, and a subtraction block,
331 which boils down to a size of $4 + 2 \times 8 + 3 = 23$ and a depth of $1 + 1 + 5 + 2 = 9$.

332 ρ_1^N and ρ_N^N (from (20)) can be constructed more simply, requiring 3 EML blocks (one of which to
333 get the constant $1/2$), one tanh block, and one subtraction block, thus corresponding to a size of
334 $3 + 8 + 3 = 14$ and depth of $1 + 1 + 5 + 2 = 9$. But, for the sake of simplicity, we will consider all
335 instances of ρ_m^N used to be those corresponding to Figure 7.

336 Then, we construct $\Phi_j^{N,d}$, using (23) by multiplying d instances of the one-dimensional partitions,
337 $\rho_{j_i}^N$. This requires $d-1$ multiplication blocks, and thus the EML tree corresponding to $\Phi_j^{N,d}$ has size
338 $23d + 6(d-1) = 29d - 6$ and depth $9 + 4(d-1) = 4d + 5$, yielding the following lemma:

339 **Lemma 7** (Construction of Approximate Partitions of Unity). *For $N, d \in \mathbb{N}$ and $j \in \mathcal{A}^N$, there exists*
340 *an EML tree T_θ such that $T_\theta = \Phi_j^N$ with*

$$N(T_\theta) \leq 29d - 6, \quad \text{and} \quad \text{Depth}(T_\theta) \leq 4d + 5. \quad (28)$$

341 G Proof of Theorem 1

342 *Proof.* We set out to perform the steps for the proof as described in the main text above and in Figure
343 1.

344 **Step 1. Approximation by Local Polynomials.** We divide the domain $[0, 1]^d$ into cubes of side
345 length $1/N$. To denote these cubes, we use the corresponding index set

$$\mathcal{A}^N = \{1, \dots, N\}^d \quad (29)$$

346 of cardinality $|\mathcal{A}^N| = N^d$. For every $j \in \mathcal{A}^N$, we associate the cube defined as

$$I_j^N = \bigotimes_{i=1}^d \left(\frac{j_i - 1}{N}, \frac{j_i}{N} \right). \quad (30)$$

347 We also define overlapping cubes, which will be used to satisfy the condition for the Bramble-Hilbert
348 lemma, for the given N and every $j \in \mathcal{A}^N$ as

$$J_j^N = \bigotimes_{i=1}^d \left(\frac{j_i - 2}{N}, \frac{j_i + 1}{N} \right), \quad (31)$$

349 which is of diameter $\frac{3\sqrt{d}}{N}$. We notice that there exists a ball of diameter $3/N$ such that J_j^N is
 350 star-shaped with respect to every point in this ball (per Definition 1). Thus, for every $j \in \mathcal{A}^N$, by the
 351 Bramble-Hilbert lemma (Lemma 2), there exists a polynomial p_j^N of degree at most $s-1$ such that,
 352 for all $0 \leq k \leq s-1$

$$\begin{aligned} \|f - p_j^N\|_{W^{k,\infty}(J_j^N)} &\leq \frac{\pi^{1/4}\sqrt{s}}{(s-k-1)!} \left(\frac{5d^2}{N}\right)^{s-k} |f|_{W^{s,\infty}([0,1]^d)} \\ &\leq \max_{0 \leq m \leq k} \frac{\pi^{1/4}\sqrt{s}(5d^2)^{s-m}}{(s-m-1)!} \frac{|f|_{W^{s,\infty}([0,1]^d)}}{N^{s-k}} := \frac{\mathcal{C}(d, k, s, f)}{N^{s-k}}, \end{aligned} \quad (32)$$

353 choosing $N > 5d^2$ (to obtain decay) and using $3\sqrt{e} \leq 5$.

354

355 **Remark 6.** *The polynomials are chosen to be good approximators on the overlapping cubes (not on*
 356 *the disjoint ones) because*

357 **Step 2. Definition of the Global Polynomial.** We then define the global piecewise polynomial over
 358 the entire domain $[0, 1]^d$ by

$$p^N = \sum_{j \in \mathcal{A}^N} p_j^N \chi_j, \quad (33)$$

359 where χ_j denotes the indicator function on I_j^N . Since p^N is composed using indicator functions
 360 (which are discontinuous), it cannot be adequately controlled in $W^{k,\infty}$ spaces where $k \geq 1$. This is
 361 why we introduce

$$\tilde{p}^N = \sum_{j \in \mathcal{A}^N} p_j^N \Phi_j^N, \quad (34)$$

362 as the global polynomial using the approximate partitions of unity and which we will ultimately
 363 represent using an EML tree. Similarly, we define $\tilde{f}^N$ as

$$\tilde{f}^N = \sum_{j \in \mathcal{A}^N} f \Phi_j^N. \quad (35)$$

364 **Step 3. Approximation of f by $\tilde{f}^N$.** We first prove that $\tilde{f}^N$ is close to f . Considering an arbitrary
 365 $i \in \mathcal{A}^N$ and any $k \geq 1$, we employ the property 1, Lemmas 5 and 6, and the explicit value of α (from
 366 (22)) to get

$$\begin{aligned} \|f - \tilde{f}^N\|_{W^{k,\infty}(I_i^N)} &= \left\| f - \sum_{j \in \mathcal{A}^N} f \Phi_j^N \right\|_{W^{k,\infty}(I_i^N)} \\ &\leq 2^k \|f\|_{W^{k,\infty}(I_i^N)} \left\| 1 - \sum_{v \in \mathcal{V}_d} \Phi_{i+v}^N \right\|_{W^{k,\infty}(I_i^N)} + 2^k \|f\|_{W^{k,\infty}(I_i^N)} \left\| \sum_{\substack{j \in \mathcal{A}^N \\ j-i \notin \mathcal{V}_d}} \Phi_j^N \right\|_{W^{k,\infty}(I_i^N)} \\ &\leq 2^k \|f\|_{W^{k,\infty}(I_i^N)} (2^{kd} d\varepsilon + N^d (2k)^{2k} \alpha^k \varepsilon) \\ &\leq 2^k \|f\|_{W^{k,\infty}(I_i^N)} \left(2^{kd} d + N^d (2k)^{2k} N^k (k+1)^k \max \left\{ R^k, \ln^k \left(\frac{2Nk^2}{\varepsilon^{\frac{1}{k+1}} e} \right) \right\} \right) \varepsilon \\ &\leq \delta (2(k+1))^{3k} \max \left\{ R^k, \ln^k \left(\frac{2Nk^2}{\varepsilon^{\frac{1}{k+1}} e} \right) \right\} \frac{\mathcal{C}(d, k, s, f)}{N^{s-k}}, \end{aligned} \quad (36)$$

367 with ε chosen to satisfy

$$\varepsilon \leq \frac{\delta \mathcal{C}(d, k, s, f)}{2^{(k+1)d} d N^{s+d} \|f\|_{W^{k,\infty}([0,1]^d)}}. \quad (37)$$

368 Similarly, for $k = 0$, we have

$$\|f - \tilde{f}^N\|_{L^\infty(I_i^N)} \leq \|f\|_{L^\infty(I_i^N)} (d\varepsilon + N^d \varepsilon) \leq \frac{\delta \mathcal{C}(d, k, s, f)}{3 N^{s-k}}. \quad (38)$$

369 **Step 4. Approximation of $\tilde{f}$ by the Global Polynomial $\tilde{p}^N$.** For $k = 0$,

$$\begin{aligned}
\|\tilde{f}^N - \tilde{p}^N\|_{L^\infty(I_i^N)} &= \left\| \sum_{j \in \mathcal{A}^N} (f - p_j^N) \Phi_j^N \right\|_{L^\infty(I_i^N)} \\
&= \left\| \sum_{v \in \mathcal{V}_d} (f - p_{i+v}^N) \Phi_{i+v}^N \right\|_{L^\infty(I_i^N)} + \sum_{\substack{j \in \mathcal{A}^N \\ j-i \notin \mathcal{V}_d}} \|f - p_j^N\|_{L^\infty(I_i^N)} \|\Phi_j^N\|_{L^\infty(I_i^N)} \\
&\leq \left(\max_{v \in \mathcal{V}_d} \|f - p_{i+v}^N\|_{L^\infty(I_i^N)} \right) \left\| \sum_{v \in \mathcal{V}_d} \Phi_{i+v}^N \right\|_{L^\infty(I_i^N)} + \frac{\mathcal{C}(d, k, s, f)}{N^{s-k}} N^d \varepsilon \\
&\leq \frac{\mathcal{C}(d, k, s, f)}{N^{s-k}} (1 + d\varepsilon) + \frac{\mathcal{C}(d, k, s, f)}{N^{s-k}} N^d \varepsilon \\
&\leq \left(1 + \frac{\delta}{3}\right) \frac{\mathcal{C}(d, k, s, f)}{N^{s-k}},
\end{aligned} \tag{39}$$

370 with suitably chosen ε .

371 Whereas for $k \geq 1$,

$$\begin{aligned}
\|\tilde{f}^N - \tilde{p}^N\|_{W^{k, \infty}(I_i^N)} &= \left\| \sum_{j \in \mathcal{A}^N} (f - p_j^N) \Phi_j^N \right\|_{W^{k, \infty}(I_i^N)} \\
&= \left\| \sum_{v \in \mathcal{V}_d} (f - p_{i+v}^N) \Phi_{i+v}^N \right\|_{W^{k, \infty}(I_i^N)} + \sum_{\substack{j \in \mathcal{A}^N \\ j-i \notin \mathcal{V}_d}} 2^k \|f - p_j^N\|_{L^\infty(I_i^N)} \|\Phi_j^N\|_{W^{k, \infty}(I_i^N)} \\
&\leq \text{something} + N^d 2^k
\end{aligned} \tag{40}$$

372 **Step 5. Finding the EML Tree.** In this step, our objective is to construct an EML tree T^N such that
373 it is exactly equal to the approximate global polynomial, that is $T^N = \tilde{p}^N$.

374 **Step 6. Conclusion.** Finally,

$$\begin{aligned}
\|f - T^N\|_{W^{k, \infty}([0,1]^d)} &\leq \|f - \tilde{f}^N\|_{W^{k, \infty}([0,1]^d)} + \|\tilde{f}^N - \tilde{p}^N\|_{W^{k, \infty}([0,1]^d)} + \|\tilde{p}^N - T^N\|_{W^{k, \infty}([0,1]^d)} \\
&\leq \|f - \tilde{f}^N\|_{W^{k, \infty}([0,1]^d)} + \|\tilde{f}^N - \tilde{p}^N\|_{W^{k, \infty}([0,1]^d)} \\
&\leq
\end{aligned} \tag{41}$$

375 **Step 7. Upper Bound on the Size of the EML Tree.**

376

□

References

- Tim De Ryck, Samuel Lanthaler, and Siddhartha Mishra. On the approximation of functions by tanh neural networks. *Neural Networks*, 143:732–750, 2021.
- Ricardo G Durán. On polynomial approximation in sobolev spaces. *SIAM journal on numerical analysis*, 20(5):985–988, 1983.
- Andrzej Odrzywołek. All elementary functions from a single binary operator, 2026. URL <https://arxiv.org/abs/2603.21852>.
- Rüdiger Verfürth. A note on polynomial approximation in sobolev spaces. *ESAIM: Mathematical Modelling and Numerical Analysis*, 33(4):715–719, 1999.
- Philippe Martin Wyder, Judah Goldfeder, Alexey Yermakov, Yue Zhao, Stefano Riva, Jan P Williams, David Zoro, Amy Sara Rude, Matteo Tomasetto, Joe Germany, et al. Common task framework for a critical evaluation of scientific machine learning algorithms. *arXiv preprint arXiv:2510.23166*, 2025.
- Alexey Yermakov, Yue Zhao, Marine Denolle, Yiyu Ni, Philippe M Wyder, Judah Goldfeder, Stefano Riva, Jan Williams, David Zoro, Amy Sara Rude, et al. The seismic wavefield common task framework. *arXiv preprint arXiv:2512.19927*, 2025.

NeurIPS Paper Checklist

The checklist is designed to encourage best practices for responsible machine learning research, addressing issues of reproducibility, transparency, research ethics, and societal impact. Do not remove the checklist: **The papers not including the checklist will be desk rejected.** The checklist should follow the references and follow the (optional) supplemental material. The checklist does NOT count towards the page limit.

Please read the checklist guidelines carefully for information on how to answer these questions. For each question in the checklist:

- You should answer [Yes], [No], or [N/A].
- [N/A] means either that the question is Not Applicable for that particular paper or the relevant information is Not Available.
- Please provide a short (1–2 sentence) justification right after your answer (even for [N/A]).

The checklist answers are an integral part of your paper submission. They are visible to the reviewers, area chairs, senior area chairs, and ethics reviewers. You will also be asked to include it (after eventual revisions) with the final version of your paper, and its final version will be published with the paper.

The reviewers of your paper will be asked to use the checklist as one of the factors in their evaluation. While [Yes] is generally preferable to [No], it is perfectly acceptable to answer [No] provided a proper justification is given (e.g., error bars are not reported because it would be too computationally expensive” or “we were unable to find the license for the dataset we used”). In general, answering [No] or [N/A] is not grounds for rejection. While the questions are phrased in a binary way, we acknowledge that the true answer is often more nuanced, so please just use your best judgment and write a justification to elaborate. All supporting evidence can appear either in the main paper or the supplemental material, provided in appendix. If you answer [Yes] to a question, in the justification please point to the section(s) where related material for the question can be found.

IMPORTANT, please:

- **Delete this instruction block, but keep the section heading “NeurIPS Paper Checklist”.**
- **Keep the checklist subsection headings, questions/answers and guidelines below.**
- **Do not modify the questions and only use the provided macros for your answers.**

1. Claims

Question: Do the main claims made in the abstract and introduction accurately reflect the paper’s contributions and scope?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the abstract and introduction do not include the claims made in the paper.
- The abstract and/or introduction should clearly state the claims made, including the contributions made in the paper and important assumptions and limitations. A [No] or [N/A] answer to this question will not be perceived well by the reviewers.
- The claims made should match theoretical and experimental results, and reflect how much the results can be expected to generalize to other settings.
- It is fine to include aspirational goals as motivation as long as it is clear that these goals are not attained by the paper.

2. Limitations

Question: Does the paper discuss the limitations of the work performed by the authors?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper has no limitation while the answer [No] means that the paper has limitations, but those are not discussed in the paper.
- The authors are encouraged to create a separate “Limitations” section in their paper.
- The paper should point out any strong assumptions and how robust the results are to violations of these assumptions (e.g., independence assumptions, noiseless settings, model well-specification, asymptotic approximations only holding locally). The authors should reflect on how these assumptions might be violated in practice and what the implications would be.
- The authors should reflect on the scope of the claims made, e.g., if the approach was only tested on a few datasets or with a few runs. In general, empirical results often depend on implicit assumptions, which should be articulated.
- The authors should reflect on the factors that influence the performance of the approach. For example, a facial recognition algorithm may perform poorly when image resolution is low or images are taken in low lighting. Or a speech-to-text system might not be used reliably to provide closed captions for online lectures because it fails to handle technical jargon.
- The authors should discuss the computational efficiency of the proposed algorithms and how they scale with dataset size.
- If applicable, the authors should discuss possible limitations of their approach to address problems of privacy and fairness.
- While the authors might fear that complete honesty about limitations might be used by reviewers as grounds for rejection, a worse outcome might be that reviewers discover limitations that aren’t acknowledged in the paper. The authors should use their best judgment and recognize that individual actions in favor of transparency play an important role in developing norms that preserve the integrity of the community. Reviewers will be specifically instructed to not penalize honesty concerning limitations.

3. Theory assumptions and proofs

Question: For each theoretical result, does the paper provide the full set of assumptions and a complete (and correct) proof?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include theoretical results.
- All the theorems, formulas, and proofs in the paper should be numbered and cross-referenced.
- All assumptions should be clearly stated or referenced in the statement of any theorems.
- The proofs can either appear in the main paper or the supplemental material, but if they appear in the supplemental material, the authors are encouraged to provide a short proof sketch to provide intuition.
- Inversely, any informal proof provided in the core of the paper should be complemented by formal proofs provided in appendix or supplemental material.
- Theorems and Lemmas that the proof relies upon should be properly referenced.

4. Experimental result reproducibility

Question: Does the paper fully disclose all the information needed to reproduce the main experimental results of the paper to the extent that it affects the main claims and/or conclusions of the paper (regardless of whether the code and data are provided or not)?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not include experiments.

- If the paper includes experiments, a [No] answer to this question will not be perceived well by the reviewers: Making the paper reproducible is important, regardless of whether the code and data are provided or not.
- If the contribution is a dataset and/or model, the authors should describe the steps taken to make their results reproducible or verifiable.
- Depending on the contribution, reproducibility can be accomplished in various ways. For example, if the contribution is a novel architecture, describing the architecture fully might suffice, or if the contribution is a specific model and empirical evaluation, it may be necessary to either make it possible for others to replicate the model with the same dataset, or provide access to the model. In general, releasing code and data is often one good way to accomplish this, but reproducibility can also be provided via detailed instructions for how to replicate the results, access to a hosted model (e.g., in the case of a large language model), releasing of a model checkpoint, or other means that are appropriate to the research performed.
- While NeurIPS does not require releasing code, the conference does require all submissions to provide some reasonable avenue for reproducibility, which may depend on the nature of the contribution. For example
 - (a) If the contribution is primarily a new algorithm, the paper should make it clear how to reproduce that algorithm.
 - (b) If the contribution is primarily a new model architecture, the paper should describe the architecture clearly and fully.
 - (c) If the contribution is a new model (e.g., a large language model), then there should either be a way to access this model for reproducing the results or a way to reproduce the model (e.g., with an open-source dataset or instructions for how to construct the dataset).
 - (d) We recognize that reproducibility may be tricky in some cases, in which case authors are welcome to describe the particular way they provide for reproducibility. In the case of closed-source models, it may be that access to the model is limited in some way (e.g., to registered users), but it should be possible for other researchers to have some path to reproducing or verifying the results.

5. Open access to data and code

Question: Does the paper provide open access to the data and code, with sufficient instructions to faithfully reproduce the main experimental results, as described in supplemental material?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that paper does not include experiments requiring code.
- Please see the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- While we encourage the release of code and data, we understand that this might not be possible, so [No] is an acceptable answer. Papers cannot be rejected simply for not including code, unless this is central to the contribution (e.g., for a new open-source benchmark).
- The instructions should contain the exact command and environment needed to run to reproduce the results. See the NeurIPS code and data submission guidelines (<https://neurips.cc/public/guides/CodeSubmissionPolicy>) for more details.
- The authors should provide instructions on data access and preparation, including how to access the raw data, preprocessed data, intermediate data, and generated data, etc.
- The authors should provide scripts to reproduce all experimental results for the new proposed method and baselines. If only a subset of experiments are reproducible, they should state which ones are omitted from the script and why.
- At submission time, to preserve anonymity, the authors should release anonymized versions (if applicable).

- Providing as much information as possible in supplemental material (appended to the paper) is recommended, but including URLs to data and code is permitted.

6. Experimental setting/details

Question: Does the paper specify all the training and test details (e.g., data splits, hyperparameters, how they were chosen, type of optimizer) necessary to understand the results?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The experimental setting should be presented in the core of the paper to a level of detail that is necessary to appreciate the results and make sense of them.
- The full details can be provided either with the code, in appendix, or as supplemental material.

7. Experiment statistical significance

Question: Does the paper report error bars suitably and correctly defined or other appropriate information about the statistical significance of the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The authors should answer **[Yes]** if the results are accompanied by error bars, confidence intervals, or statistical significance tests, at least for the experiments that support the main claims of the paper.
- The factors of variability that the error bars are capturing should be clearly stated (for example, train/test split, initialization, random drawing of some parameter, or overall run with given experimental conditions).
- The method for calculating the error bars should be explained (closed form formula, call to a library function, bootstrap, etc.)
- The assumptions made should be given (e.g., Normally distributed errors).
- It should be clear whether the error bar is the standard deviation or the standard error of the mean.
- It is OK to report 1-sigma error bars, but one should state it. The authors should preferably report a 2-sigma error bar than state that they have a 96% CI, if the hypothesis of Normality of errors is not verified.
- For asymmetric distributions, the authors should be careful not to show in tables or figures symmetric error bars that would yield results that are out of range (e.g., negative error rates).
- If error bars are reported in tables or plots, the authors should explain in the text how they were calculated and reference the corresponding figures or tables in the text.

8. Experiments compute resources

Question: For each experiment, does the paper provide sufficient information on the computer resources (type of compute workers, memory, time of execution) needed to reproduce the experiments?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer **[N/A]** means that the paper does not include experiments.
- The paper should indicate the type of compute workers CPU or GPU, internal cluster, or cloud provider, including relevant memory and storage.
- The paper should provide the amount of compute required for each of the individual experimental runs as well as estimate the total compute.

- 598 • The paper should disclose whether the full research project required more compute
599 than the experiments reported in the paper (e.g., preliminary or failed experiments that
600 didn't make it into the paper).

601 9. Code of ethics

602 Question: Does the research conducted in the paper conform, in every respect, with the
603 NeurIPS Code of Ethics <https://neurips.cc/public/EthicsGuidelines?>

604 Answer: **[TODO]**

605 Justification: **[TODO]**

606 Guidelines:

- 607 • The answer [N/A] means that the authors have not reviewed the NeurIPS Code of
608 Ethics.
- 609 • If the authors answer **[No]**, they should explain the special circumstances that require a
610 deviation from the Code of Ethics.
- 611 • The authors should make sure to preserve anonymity (e.g., if there is a special consid-
612 eration due to laws or regulations in their jurisdiction).

613 10. Broader impacts

614 Question: Does the paper discuss both potential positive societal impacts and negative
615 societal impacts of the work performed?

616 Answer: **[TODO]**

617 Justification: **[TODO]**

618 Guidelines:

- 619 • The answer [N/A] means that there is no societal impact of the work performed.
- 620 • If the authors answer [N/A] or **[No]**, they should explain why their work has no societal
621 impact or why the paper does not address societal impact.
- 622 • Examples of negative societal impacts include potential malicious or unintended uses
623 (e.g., disinformation, generating fake profiles, surveillance), fairness considerations
624 (e.g., deployment of technologies that could make decisions that unfairly impact specific
625 groups), privacy considerations, and security considerations.
- 626 • The conference expects that many papers will be foundational research and not tied
627 to particular applications, let alone deployments. However, if there is a direct path to
628 any negative applications, the authors should point it out. For example, it is legitimate
629 to point out that an improvement in the quality of generative models could be used to
630 generate Deepfakes for disinformation. On the other hand, it is not needed to point out
631 that a generic algorithm for optimizing neural networks could enable people to train
632 models that generate Deepfakes faster.
- 633 • The authors should consider possible harms that could arise when the technology is
634 being used as intended and functioning correctly, harms that could arise when the
635 technology is being used as intended but gives incorrect results, and harms following
636 from (intentional or unintentional) misuse of the technology.
- 637 • If there are negative societal impacts, the authors could also discuss possible mitigation
638 strategies (e.g., gated release of models, providing defenses in addition to attacks,
639 mechanisms for monitoring misuse, mechanisms to monitor how a system learns from
640 feedback over time, improving the efficiency and accessibility of ML).

641 11. Safeguards

642 Question: Does the paper describe safeguards that have been put in place for responsible
643 release of data or models that have a high risk for misuse (e.g., pre-trained language models,
644 image generators, or scraped datasets)?

645 Answer: **[TODO]**

646 Justification: **[TODO]**

647 Guidelines:

- 648 • The answer [N/A] means that the paper poses no such risks.

- Released models that have a high risk for misuse or dual-use should be released with necessary safeguards to allow for controlled use of the model, for example by requiring that users adhere to usage guidelines or restrictions to access the model or implementing safety filters.
- Datasets that have been scraped from the Internet could pose safety risks. The authors should describe how they avoided releasing unsafe images.
- We recognize that providing effective safeguards is challenging, and many papers do not require this, but we encourage authors to take this into account and make a best faith effort.

12. Licenses for existing assets

Question: Are the creators or original owners of assets (e.g., code, data, models), used in the paper, properly credited and are the license and terms of use explicitly mentioned and properly respected?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not use existing assets.
- The authors should cite the original paper that produced the code package or dataset.
- The authors should state which version of the asset is used and, if possible, include a URL.
- The name of the license (e.g., CC-BY 4.0) should be included for each asset.
- For scraped data from a particular source (e.g., website), the copyright and terms of service of that source should be provided.
- If assets are released, the license, copyright information, and terms of use in the package should be provided. For popular datasets, paperswithcode.com/datasets has curated licenses for some datasets. Their licensing guide can help determine the license of a dataset.
- For existing datasets that are re-packaged, both the original license and the license of the derived asset (if it has changed) should be provided.
- If this information is not available online, the authors are encouraged to reach out to the asset's creators.

13. New assets

Question: Are new assets introduced in the paper well documented and is the documentation provided alongside the assets?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not release new assets.
- Researchers should communicate the details of the dataset/code/model as part of their submissions via structured templates. This includes details about training, license, limitations, etc.
- The paper should discuss whether and how consent was obtained from people whose asset is used.
- At submission time, remember to anonymize your assets (if applicable). You can either create an anonymized URL or include an anonymized zip file.

14. Crowdsourcing and research with human subjects

Question: For crowdsourcing experiments and research with human subjects, does the paper include the full text of instructions given to participants and screenshots, if applicable, as well as details about compensation (if any)?

Answer: **[TODO]**

Justification: **[TODO]**

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Including this information in the supplemental material is fine, but if the main contribution of the paper involves human subjects, then as much detail as possible should be included in the main paper.
- According to the NeurIPS Code of Ethics, workers involved in data collection, curation, or other labor should be paid at least the minimum wage in the country of the data collector.

15. Institutional review board (IRB) approvals or equivalent for research with human subjects

Question: Does the paper describe potential risks incurred by study participants, whether such risks were disclosed to the subjects, and whether Institutional Review Board (IRB) approvals (or an equivalent approval/review based on the requirements of your country or institution) were obtained?

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the paper does not involve crowdsourcing nor research with human subjects.
- Depending on the country in which research is conducted, IRB approval (or equivalent) may be required for any human subjects research. If you obtained IRB approval, you should clearly state this in the paper.
- We recognize that the procedures for this may vary significantly between institutions and locations, and we expect authors to adhere to the NeurIPS Code of Ethics and the guidelines for their institution.
- For initial submissions, do not include any information that would break anonymity (if applicable), such as the institution conducting the review.

16. Declaration of LLM usage

Question: Does the paper describe the usage of LLMs if it is an important, original, or non-standard component of the core methods in this research? Note that if the LLM is used only for writing, editing, or formatting purposes and does *not* impact the core methodology, scientific rigor, or originality of the research, declaration is not required.

Answer: [TODO]

Justification: [TODO]

Guidelines:

- The answer [N/A] means that the core method development in this research does not involve LLMs as any important, original, or non-standard components.
- Please refer to our LLM policy in the NeurIPS handbook for what should or should not be described.